



University of  
BRISTOL

# Reproducible data science

Combining R Studio, Git & R Markdown for reproducibility

Henry W J Reeve

[henry.reeve@bristol.ac.uk](mailto:henry.reeve@bristol.ac.uk)

Statistical Computing & Empirical Methods (EMATM0061)

MSc in Data Science, Teaching block 1, 2021.

# What will we cover today?

- We will discuss the three main goals of scientific research
- We will discuss the central role of replicability in science.
- We consider the “replication crisis”.
- We will understand the importance of reproducible data analysis.
- We will introduce several tools for facilitating reproducible data analysis in R:
  - R Markdown
  - R Projects
  - Git integration.

# Three goals for scientific research


Data Science is the science of extracting information, insight and understanding from data!

**Reliable:** Scientific research should be potentially replicable by other scientists.

**Valid:** Scientific research should use principled approaches to ensure that the results of the investigation actually support the conclusions drawn.

**Importance:** Scientific research should address important questions.

# The importance of replication



Galileo Galilei, circa 1590:

Two objects of the same shape and material but different weights will fall at the same rate.

Replicated many times and in many locations, including on the moon – Apollo 15, 1971.

# The replication crisis!

- Scientific truths should be robust to repeated replications of the same experiment.
- John Ioannidis has argued that “most published research findings are false” Nature, 2005.
- Replication studies conducted by pharmaceutical companies:

Prinz, Schlange, Asadullah from Bayer Healthcare, 2011, 67 replication projects:

“In almost two-thirds of the projects, there were inconsistencies between published data and in-house”

Begley & Ellis from Amgen, 2012, 53 landmark studies,

“scientific findings were confirmed in only 6 cases”.



# Replicability vs. Reproducibility

- Scientific truths should be robust to repeated replications of the same experiment.
- **Replicability:** Different experimenters will yield the same results from different data, when an experiment is repeated under similar conditions.
- **Reproducibility:** Different scientists will yield the same results by repeating the analysis on the same data.
- Surprisingly, reproducibility is still a challenge!  
Difficult to reproduce analysis spread across a poorly organized amalgam of code & spread sheets.

Now take a break!



# Literate programming & reproducibility

- Donald Knuth emphasized the importance of literate programming:

"Instead of imagining that our main task is to instruct a computer what to do, let us concentrate rather on explaining to human beings what we want a computer to do."

- Make your code as readable as possible – both for others and your future self!
  - Include plenty of clear comments
  - Adopt sensible naming conventions
  - Aim for a simple organizational structure with succinct functions.

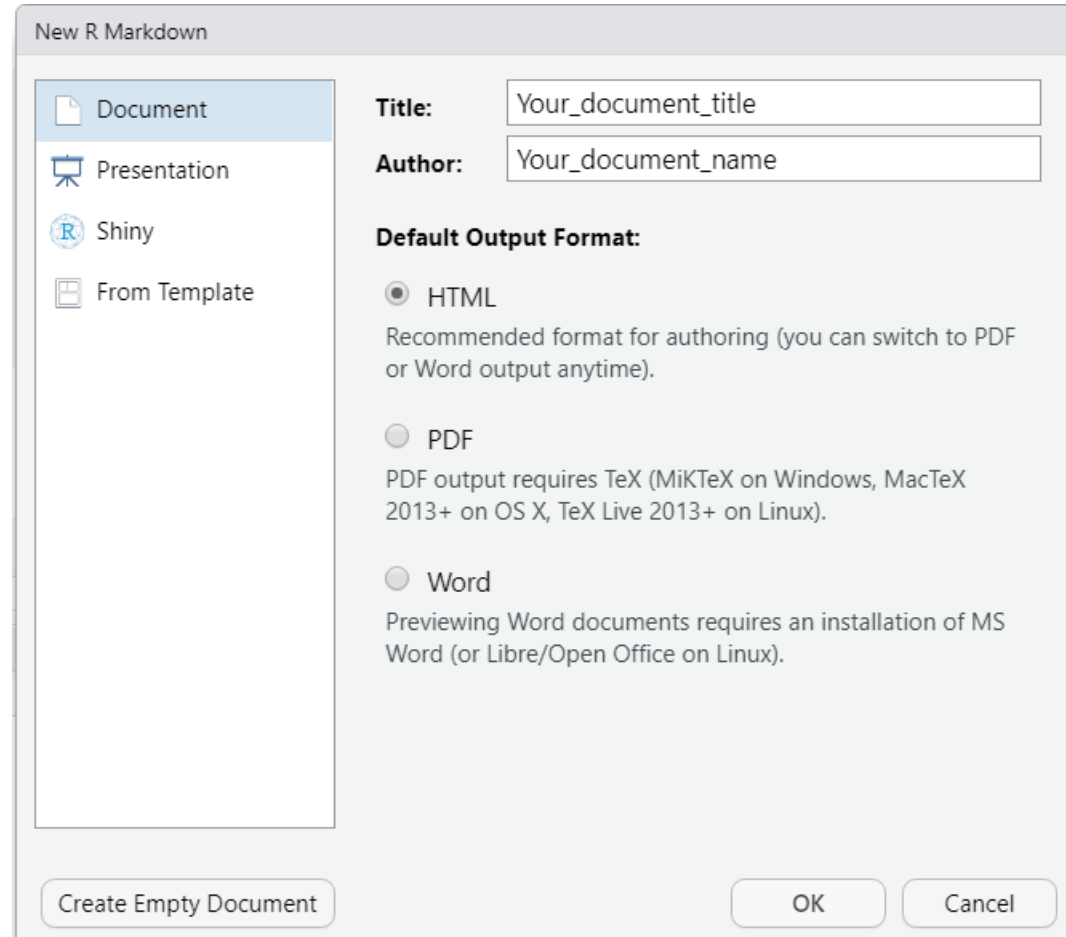


# Reproducibility with R & RStudio

- RStudio facilitates reproducible analysis via R Projects, R Markdown and Git interface.
- RMarkdown allows us to generate a notebook style document which includes R code, plots and explanatory text in a linear format.
- R Projects provide a specific workspace for each project with a working directory, data & history.
- Git is a version control system which allows us to track and revert changes, and collaborate.

# R Markdown

- We can create a new R Markdown document by File -> New File -> R Markdown ...



# R Markdown

- We can create a new R Markdown document by File -> New File -> R Markdown ...

- We can edit the title, author, date, output:

```
---  
title: "Example RMarkdown document"  
author: "Henry"  
date: "11/08/2020"  
output: html_document  
---
```

- Section headings are generated by `## R Markdown`

- We can also embed code fragments:

```
```{r building a function and a data frame}  
  
# First we create a simple function  
f <- function(z) {  
  return(5*z^2+z*cos(15*z)+0.3*sin(300*z))  
}  
  
# We randomly generate some x  
x<-runif(100)  
# We set y to be f applied to x  
y<-f(x)  
  
# We then put x and y together in a data frame  
df<-data.frame(x,y)  
```
```

- We can also include plots:

```
```{r, echo=FALSE}  
  
# A simple plot  
plot(x,y)  
```
```

# R Markdown

```
1 ---
2 title: "Example RMarkdown document"
3 output: html_document
4 ---
5
6
7 ```{r setup, include=FALSE}
8 knitr::opts_chunk$set(echo = TRUE) #We tell R to display code by default.
9 ```
10
11
12
13 ## Code fragment
14
15 We can embed pieces of R code as follows:
16
17 ```{r building a function and a data frame}
18
19 # First we create a simple function
20 f <- function(z) {
21   return(5*z^2+z*cos(15*z)+0.3*sin(300*z))
22 }
23
24 # We randomly generate some x
25 x<-runif(100)
26 # We set y to be f applied to x
27 y<-f(x)
28
29 # We then put x and y together in a data frame
30 df<-data.frame(x,y)
31
32 ```
33
34
35 ## Embedded plot
36
37 We can also embed plots. By using `echo = FALSE` we display only the output
38
39 ```{r, echo=FALSE}
40
41 # A simple plot
42 plot(x,y)
43
44 ```
```



## Example RMarkdown document

### Code fragment

We can embed pieces of R code as follows:


```
# First we create a simple function
f <- function(z) {
  return(5*z^2+z*cos(15*z)+0.3*sin(300*z))
}

# We randomly generate some x
x<-runif(100)
# We set y to be f applied to x
y<-f(x)

# We then put x and y together in a data frame
df<-data.frame(x,y)
```

### Embedded plot

We can also embed plots. By using `echo = FALSE` we display only the output and not the code.



Now take a break!






# Version control with Git

- Go to <https://github.com> and register for free GitHub account.
- Install git locally:

|                |   |
|----------------|---|
| Windows:       | <a href="https://gitforwindows.org">https://gitforwindows.org</a> |
| Mac OSX:       | <code>xcode-select --install</code>                               |
| Ubuntu/Debian: | <code>sudo apt-get install git</code>                             |
| Fedora/Redhat: | <code>sudo yum install git</code>                                 |
- Connect to your Git account within R:


```
install.packages("usethis")  
library(usethis)  
use_git_config(user.name = "Bob Smith", user.email = "bob@example.org")
```





# Set up an R project with Git version control

- Go to <https://github.com> & create a new repository by pressing  .  
Add an informative title, a description and include a README.
- Within the github repo click on  and copy the repo URL.
- Create a new project within R Studio:  
File -> New Project -> Version Control -> Git -> Enter repo URL + Project name.  
Check “Open in new session” and then create the project  .
- We can now add files, commit, push and pull using the Git panel in the top right of RStudio.

# Set up an R project with Git version control

- We can now add files, commit, push and pull using the Git panel in the top right of RStudio.



- **Stage files:** Choose which files to include in the version history by clicking  by the file name.
- **Commit:** Take a snapshot of staged files within the local git repository by  **Commit** .  
Remember to include a succinct but informative commit message.
- **Push:** Send your local changes to the master branch  .
- **Pull:** Copy changes made by your collaborators onto your local repository  .
- An excellent resource for more information from Jenny Bryan : <https://happygitwithr.com>



# What have we covered?

- We discussed the central role of replicability in science.
- We considered the idea of a “replication crisis”.
- We discussed the difference between replicability and reproducibility.
- We introduced RMarkdown for reproducible data analysis.
- We discussed the integration of Git with RStudio and R projects.



Thanks for listening,  
... now onto the assignment!

[henry.reeve@bristol.ac.uk](mailto:henry.reeve@bristol.ac.uk)

Include EMATM0061 in the subject of your email.