



An introduction to classification

Learning functions which map feature vectors to class labels

Henry W J Reeve

henry.reeve@bristol.ac.uk

Include EMATM0061 in the subject of your email.

Statistical Computing & Empirical Methods (EMATM0061)

MSc in Data Science, Teaching block 1, 2021.

What will we cover today?


- We will begin by introducing the concept of classification.
- We will emphasize the importance of predictive performance on unseen data.
- We will consider the supervised learning pipeline and the role of the test-train split.
- We will use probabilistic ideas to understand the classification problem.
- We will emphasize the difference the difference between train and test error.
- We will also discuss the fundamental concept of a Bayes classifier.

What is classification?

Learning a function $\phi : \mathcal{X} \rightarrow \mathcal{Y}$

which takes as input a feature vector $X \in \mathcal{X}$

and returns a categorical variable $\phi(X) \in \mathcal{Y}$ also known as a label.



What is classification?


Example 1: Sentiment analysis

A company wants to automatically classify social media posts as being either “positive” or “negative” in sentiment.

“The food was fantastic! ”




“The service was disappointing”



What is classification?

Example 2: Image classification


A biologist wants to automatically classify images of fish according to which species they belong to.



What is classification?

Example 3: Digit recognition



The postal service needs an automatic system for converting hand-written addresses into digitized addresses.



What is classification?

Example 3: Automated medical diagnosis

A medical doctor wants to establish an automated procedure for classifying retinal blood vessels as normal or abnormal.




What is classification?

Learning a function $\phi : \mathcal{X} \rightarrow \mathcal{Y}$ known as a **classification rule**.


which takes as input a feature vector $X \in \mathcal{X}$

and returns a categorical variable $\phi(X) \in \mathcal{Y}$



What is binary classification?

Learning a function $\phi : \mathcal{X} \rightarrow \mathcal{Y}$ known as a **classification rule**.



A binary classification problem is one with just two possible outcomes

Example: In automated medical diagnosis “normal” vs. “abnormal”.

How can we create a classification rule?

Let's suppose we want **classification rule** implemented within a computer.

The rule-based approach

We could attempt to program a detailed set of rules:

e.g. "The rainbow shark has two large eyes"...


Problems:

This approach would be incredibly labor intensive.

New problems would require new rules.

Performs poorly in practice.

Brittle e.g. what if we can't see both eyes etc.



How can we create a classification rule?

Let's suppose we want **classification rule** implemented within a computer.


The rule-based approach

We could attempt to program a detailed set of rules:

Problem: Highly labor intensive

The statistical learning approach

We set instead we design learning algorithms so that the machine can learn to classify from data.




How can we create a classification rule?

The statistical learning approach is to program our machine to learn tasks from data.

ML proverb

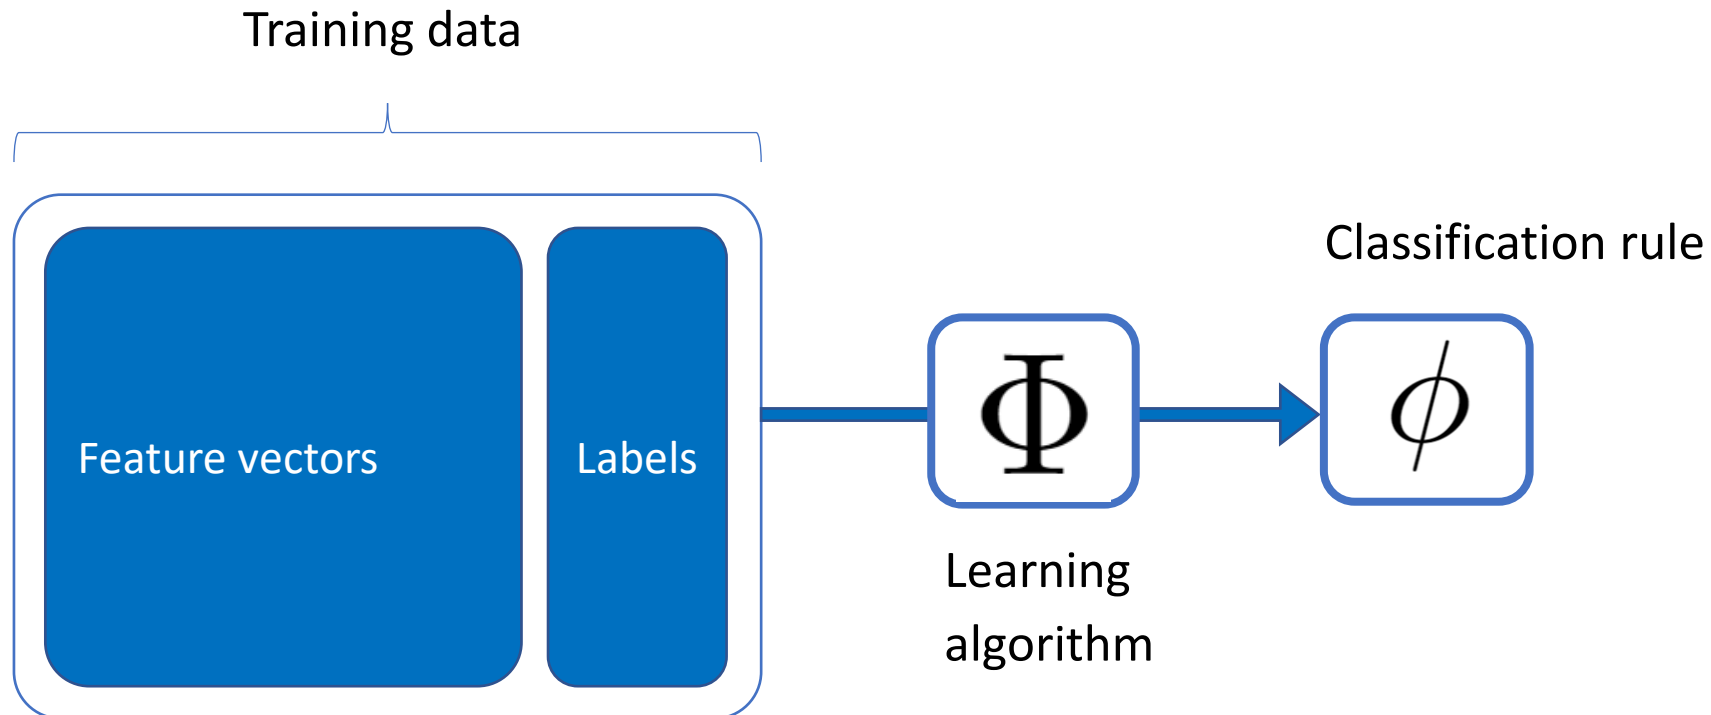
Why teach a computer to classify fish, when you can teach a computer to learn how to classify fish?



Supervised learning

We learn a **classification rule** based on a set of training data \mathcal{D} .


Training data is passed to a **learning algorithm** which outputs a classification rule.



Supervised learning

We learn a set of **classification rule** based on a set of training data \mathcal{D} .

Training data



Training data consists of a set of labelled data

$$\mathcal{D} = ((X_1, Y_1), \cdots, (X_n, Y_n))$$

A sequence of ordered pairs (X_i, Y_i) .


X_i is a feature vector.

Y_i is a label associated with X_i .

Supervised learning

We learn a set of **classification rule** based on a set of training data \mathcal{D} .

Training data



Training data consists of a set of labelled data

$$\mathcal{D} = ((X_1, Y_1), \cdots, (X_n, Y_n))$$

Example


X_i is an image of a particular fish.

Y_i is a label corresponding to the species of the fish.

Supervised learning

We learn a **classification rule** based on a set of training data \mathcal{D} .

Training data is passed to a **learning algorithm**.

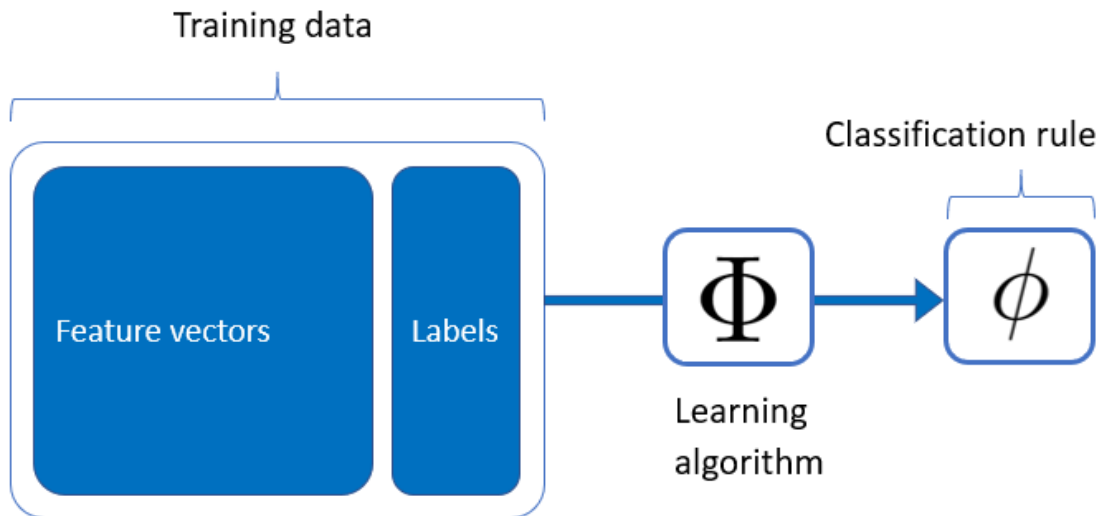


The learning algorithm
automatically identifies patterns
within the training data \mathcal{D} .

Supervised learning

We learn a set of **classification rule** based on a set of training data \mathcal{D} .

Training data is passed to a **learning algorithm** which outputs a classification rule.



The classification rule is a mapping:

$$\phi : X \mapsto Y$$

A classification rule is also known as a classifier.

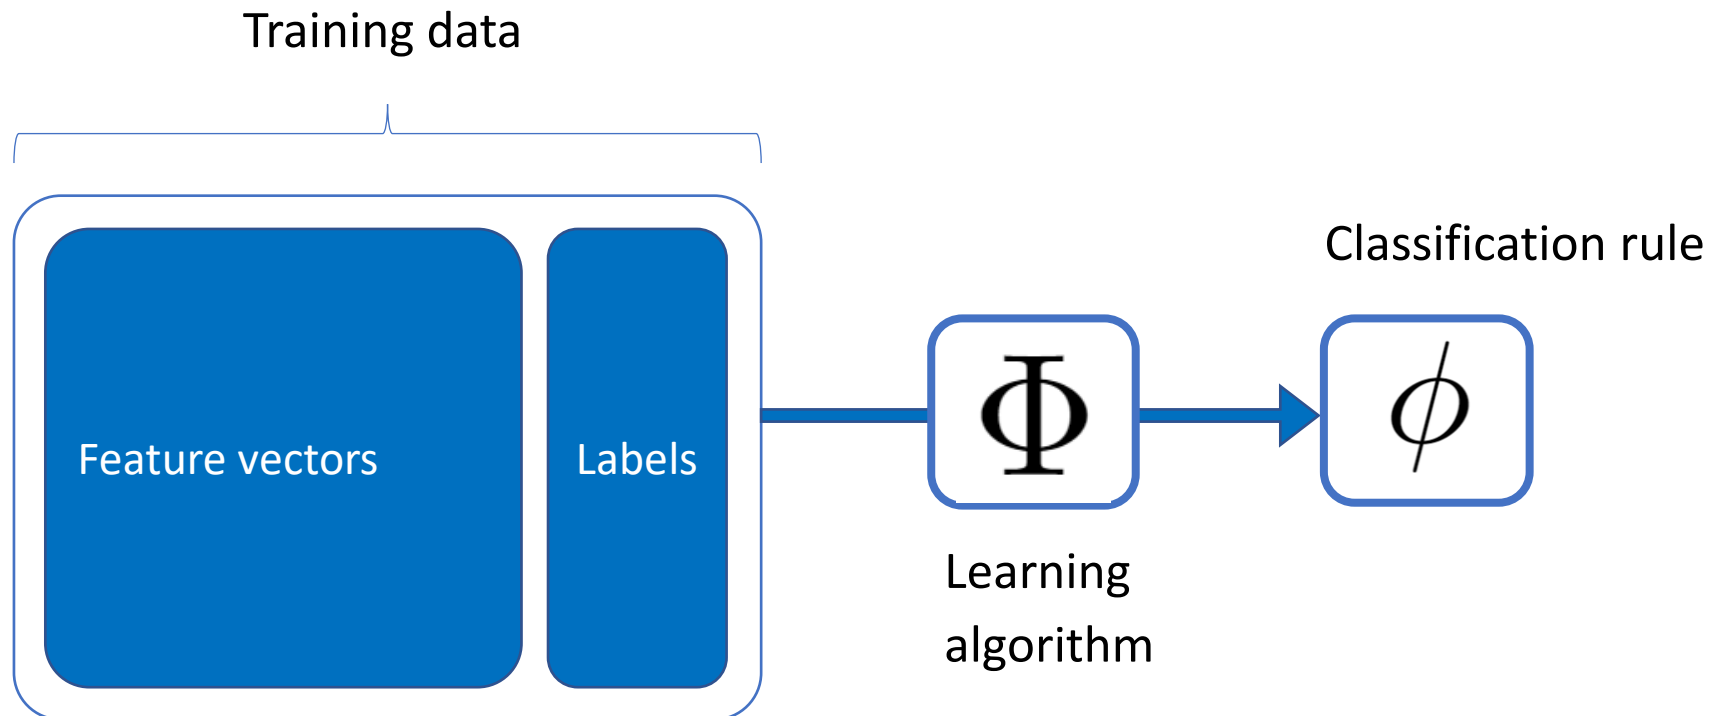
Should reflect the structure of the training data:

$$\mathcal{D} = ((X_1, Y_1), \dots, (X_n, Y_n))$$

Supervised learning

We learn a **classification rule** based on a set of training data \mathcal{D} .

Training data is passed to a **learning algorithm** which outputs a classification rule.



Now take a break!



Statistical Computing & Empirical Methods

Learning vs. memorization

Our goal is to learn a classification rule $\phi : \mathcal{X} \rightarrow \mathcal{Y}$.

The classification rule ϕ should map feature vectors $X \in \mathcal{X}$ to labels $\phi(X) \in \mathcal{Y}$.

Key point: The classification rule should perform well on unseen feature vectors $X \in \mathcal{X}$.

Not just on the training data $\mathcal{D} = ((X_1, Y_1), \dots, (X_n, Y_n))$.

Example:

We want the fish classification rule to correctly determine the fish species for new images...

... not just the images within the training data.

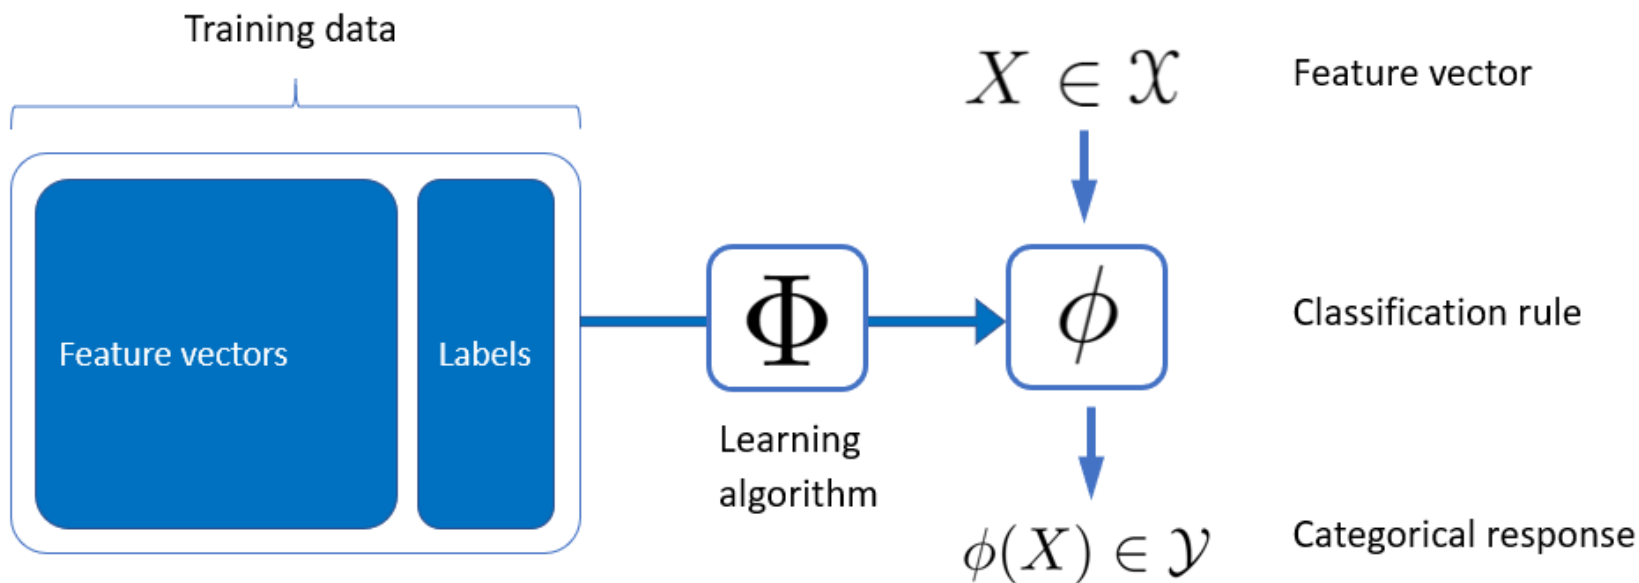
This is the crucial difference between learning and memorization.



Supervised learning

We learn a set of **classification rule** based on a set of training data \mathcal{D} .

Training data is passed to a **learning algorithm** which outputs a classification rule.

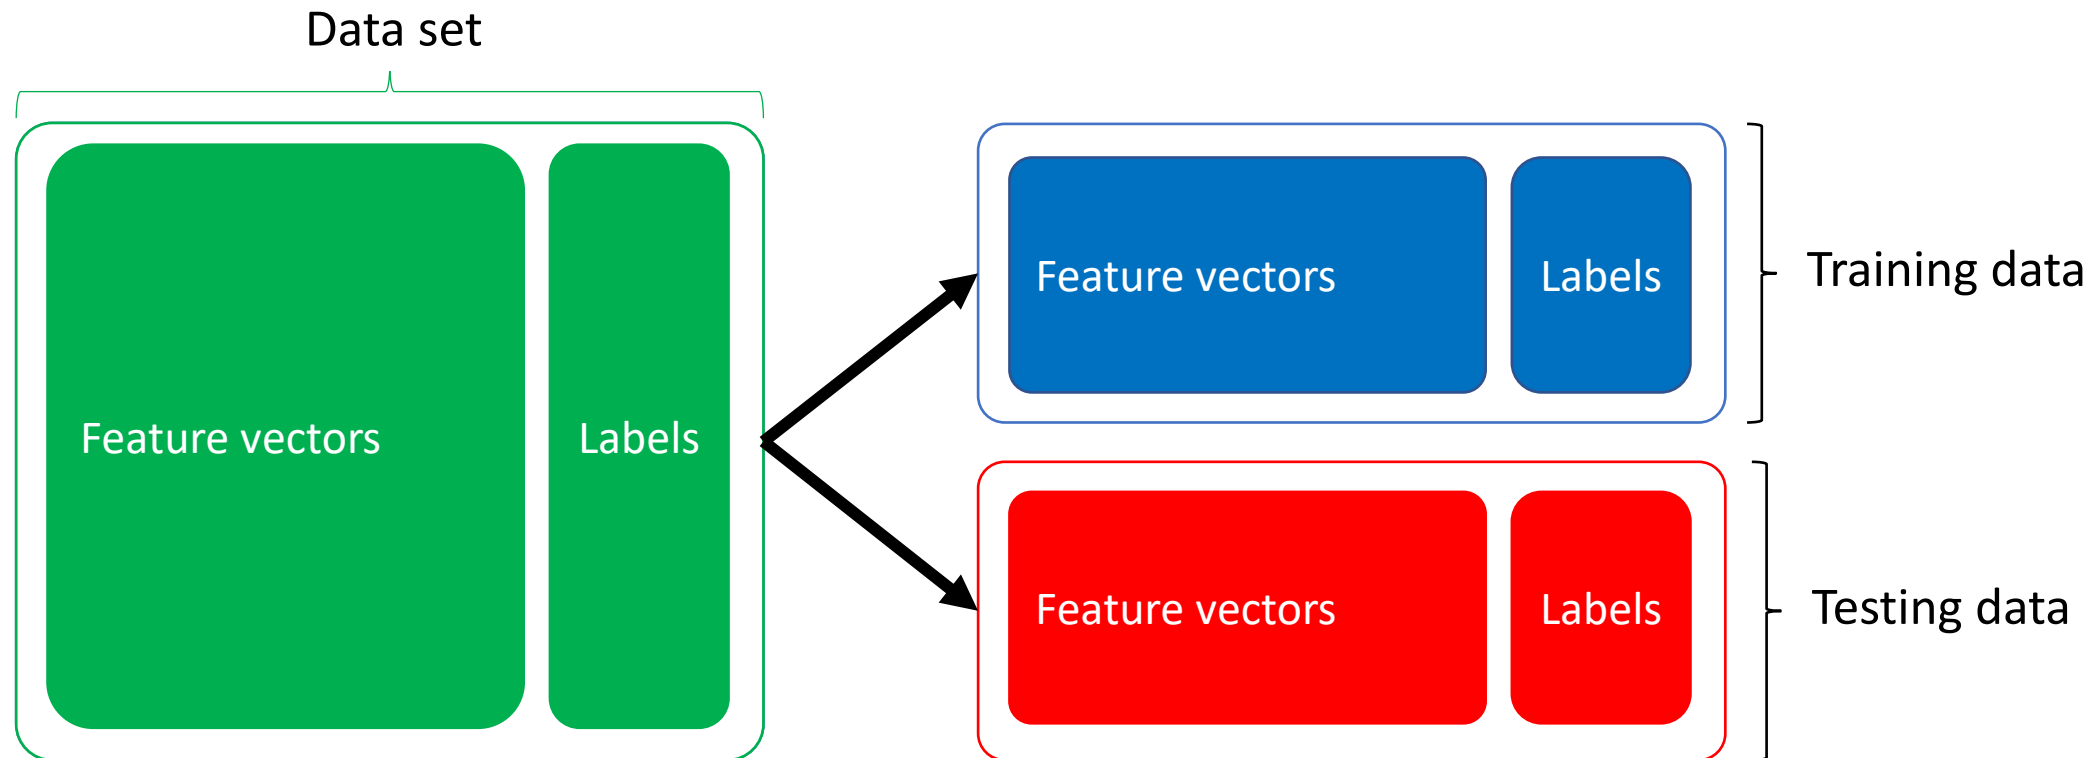


For the classification rule to be successful it must perform well on **unseen** data.

The train test split

For the classification rule to be successful it must perform well on **unseen** data.

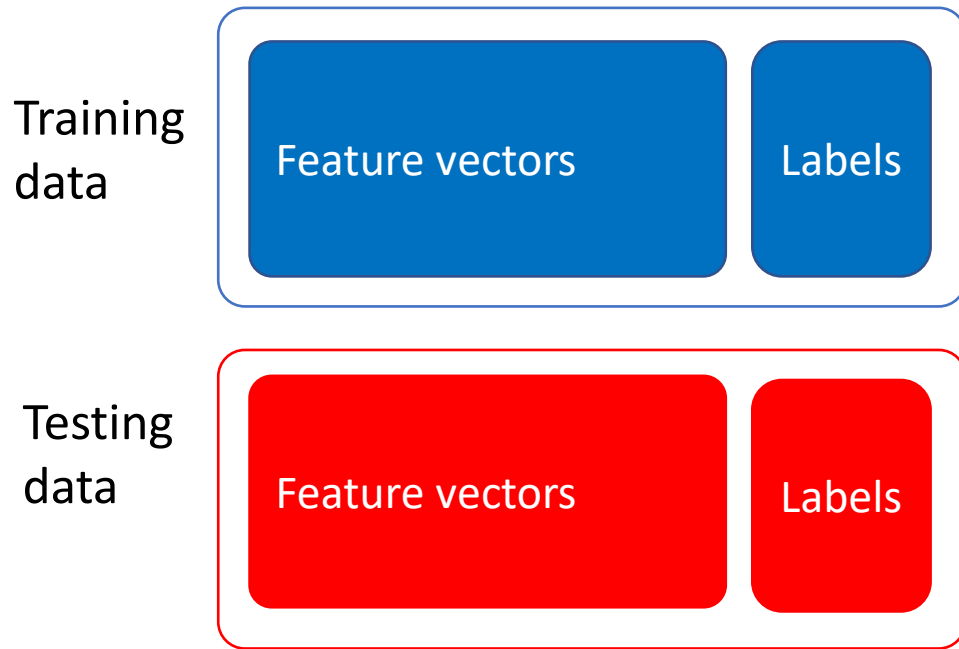
In order to investigate learning algorithms, we always need to do a test train split.



The train test split

For the classification rule to be successful it must perform well on **unseen** data.

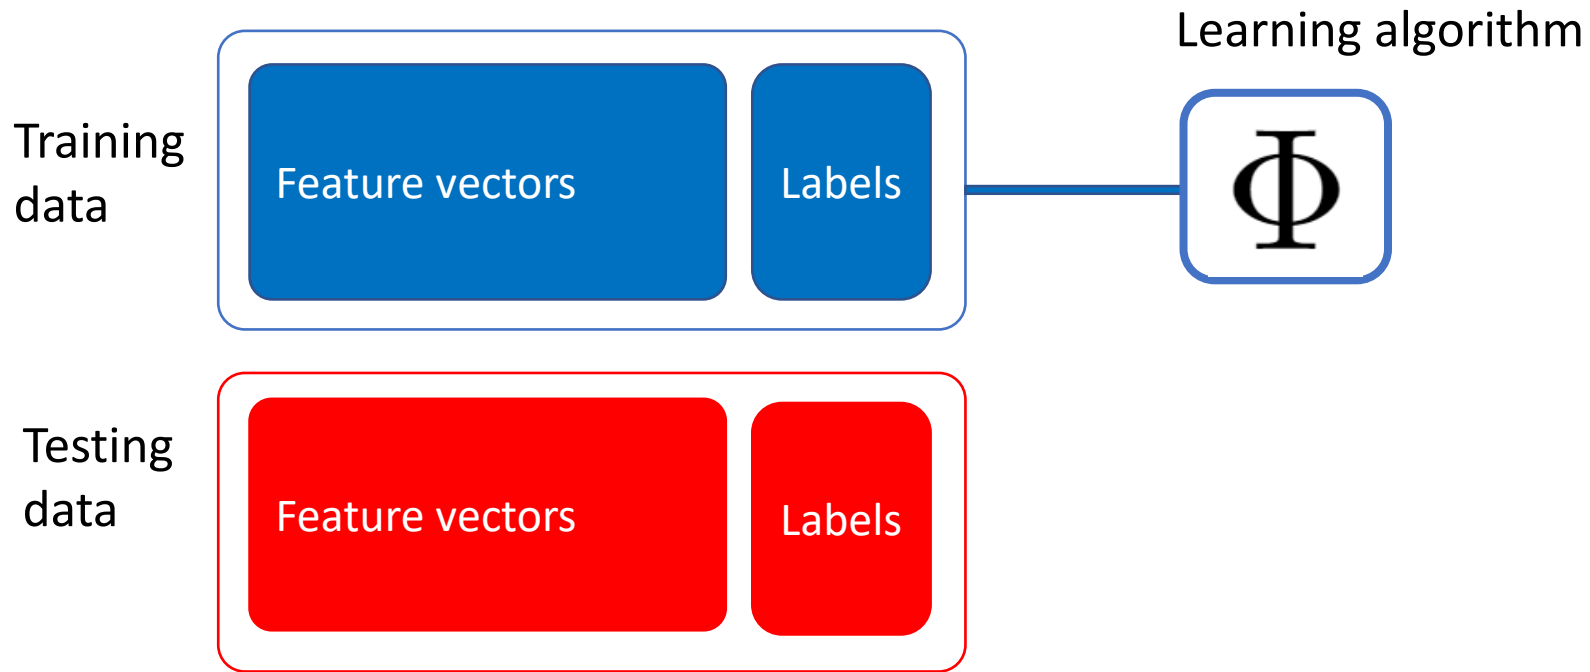
We use the train test split to assess our learning algorithms.



The train test split

For the classification rule to be successful it must perform well on **unseen** data.

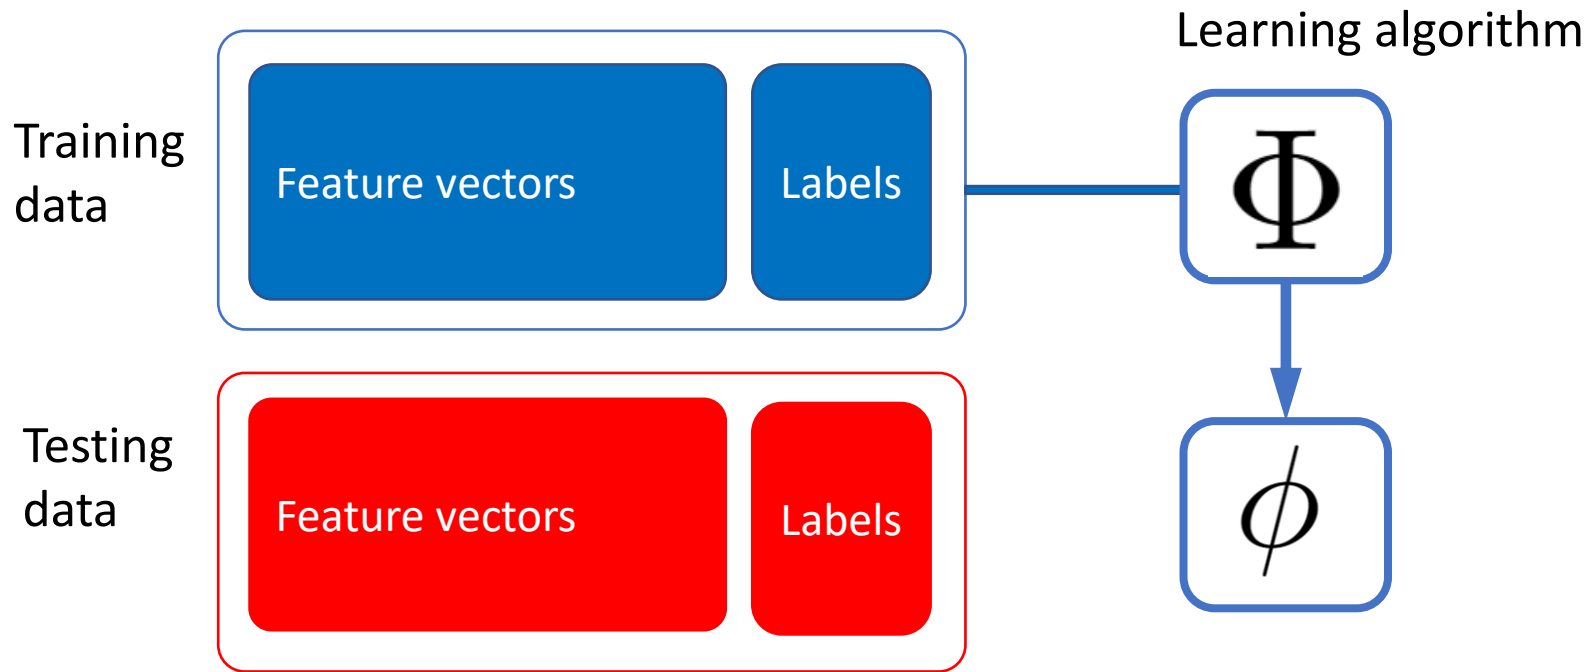
We use the train test split to assess our learning algorithms.



The train test split

For the classification rule to be successful it must perform well on **unseen** data.

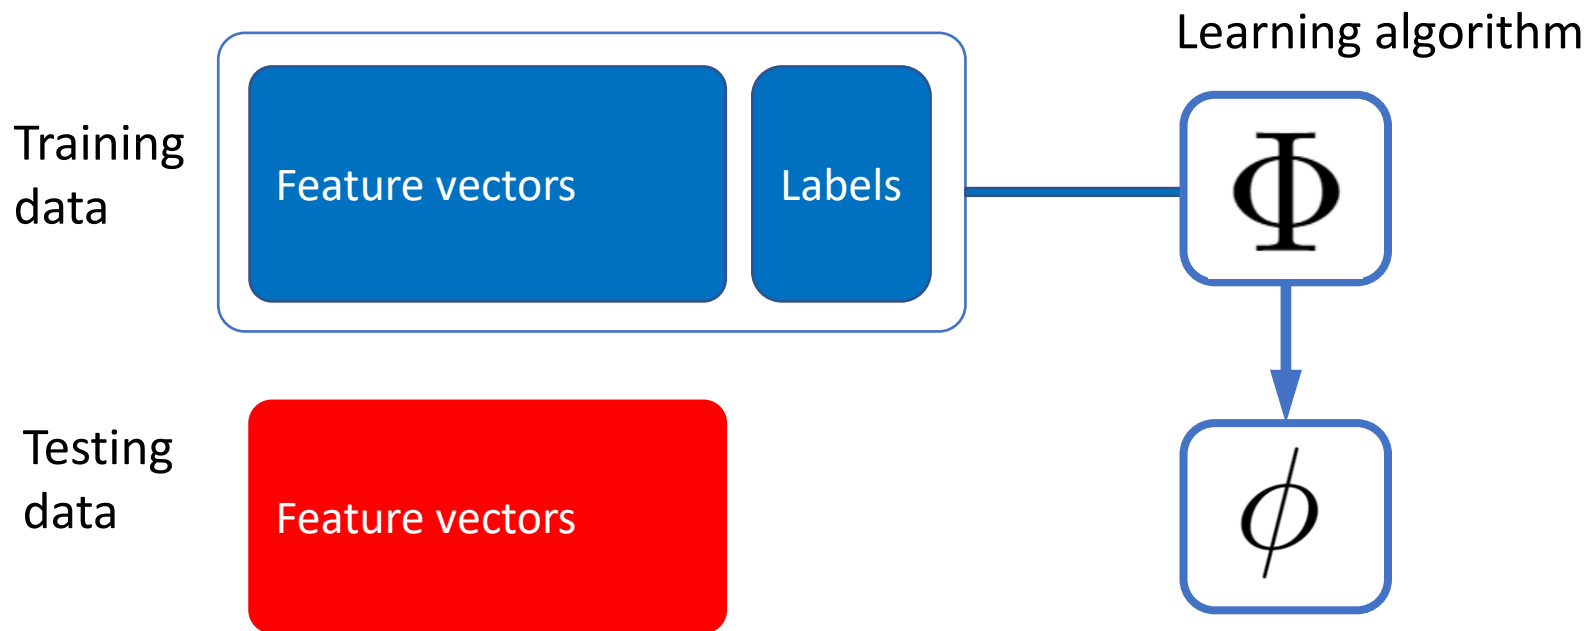
We use the train test split to assess our learning algorithms.



The train test split

For the classification rule to be successful it must perform well on **unseen** data.

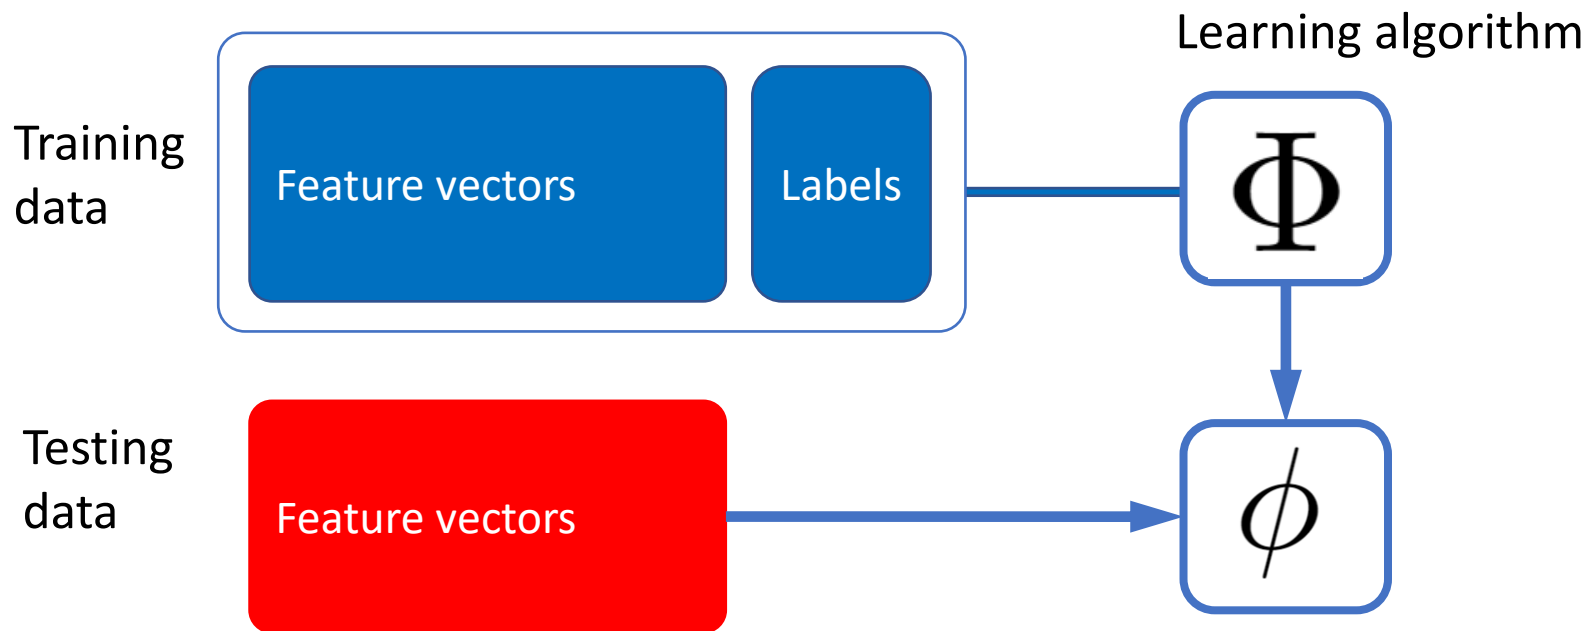
We use the train test split to assess our learning algorithms.



The train test split

For the classification rule to be successful it must perform well on **unseen** data.


We use the train test split to assess our learning algorithms.



The train test split

For the classification rule to be successful it must perform well on **unseen** data.

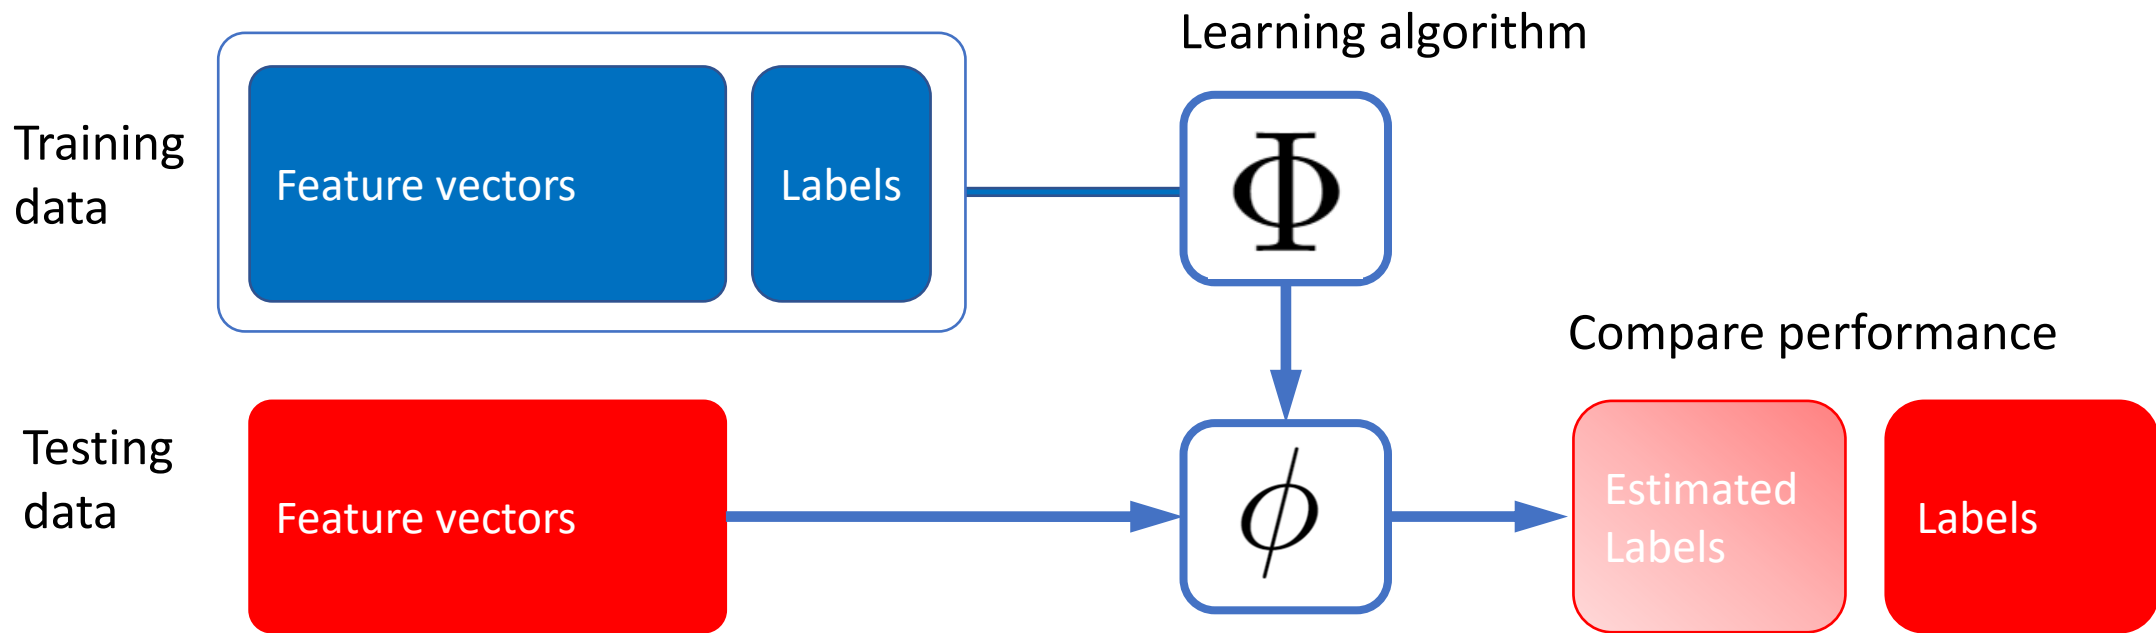
We use the train test split to assess our learning algorithms.



The train test split

For the classification rule to be successful it must perform well on **unseen** data.


We use the train test split to assess our learning algorithms.



The train test split

Key point: Never use your test data to learn your classifier!

The goal of the test data is to see how well the classification rule does on unseen data.



Example: Penguin classification

Suppose we want to learn a classifier $\phi : \mathcal{X} \rightarrow \mathcal{Y}$ which takes a feature vector of morphological features and predicts whether a penguin belongs to either the Adelie species or the Chinstrap species.

Features: $X = (X^1, X^2) \in \mathcal{X} = \mathbb{R}^2$

$X^1 =$ the weight of the penguin (grams).

$X^2 =$ the flipper length of the penguin (mm).

Labels: $Y \in \mathcal{Y} = \{0, 1\}$

$Y = \begin{cases} 1 & \text{if the penguin is an Adelie} \\ 0 & \text{if the penguin is a Chinstrap.} \end{cases}$



Example: Penguin classification

Suppose we want to learn a classifier $\phi : \mathcal{X} \rightarrow \mathcal{Y}$ which takes a feature vector of morphological features and predicts whether a penguin belongs to either the Adelie species or the Chinstrap species.

```
library(tidyverse)
library(palmerpenguins)

peng_total<-penguins%>% # prepare our data
  select(body_mass_g,flipper_length_mm,species)%>%
  filter(species!="Gentoo")%>%
  drop_na()%>%
  mutate(species=as.numeric(species=="Adelie"))
```


Example: Penguin classification

Suppose we want to learn a classifier $\phi : \mathcal{X} \rightarrow \mathcal{Y}$ which takes a feature vector of morphological features and predicts whether a penguin belongs to either the Adelie species or the Chinstrap species.

```
library(tidyverse)
library(palmerpenguins)

peng_total <- penguins %>% # prepare our data
  select(body_mass_g, flipper_length_mm, species) %>%
  filter(species != "Gentoo") %>%
  drop_na() %>%
  mutate(species = as.numeric(species == "Adelie"))
```



peng_total

```
## # A tibble: 219 x 3
##   body_mass_g flipper_length_mm species
##   <int>         <int>     <dbl>
## 1      3750           181         1
## 2      3800           186         1
## 3      3250           195         1
## 4      3450           193         1
## 5      3650           190         1
## 6      3625           181         1
## 7      4675           195         1
## 8      3475           193         1
## 9      4250           190         1
## 10     3300           186         1
## # ... with 209 more rows
```

Example: Penguin classification

Suppose we want to learn a classifier $\phi : \mathcal{X} \rightarrow \mathcal{Y}$ which takes a feature vector of morphological features and predicts whether a penguin belongs to either the Adelie species or the Chinstrap species

X Y

##	#	A tibble: 219 x 3		
##		body_mass_g	flipper_length_mm	species
##		<int>	<int>	<dbl>
##	1	3750	181	1
##	2	3800	186	1
##	3	3250	195	1
##	4	3450	193	1
##	5	3650	190	1

Feature vector $X = (X^1, X^2) \in \mathcal{X} = \mathbb{R}^2$

$X^1 =$ the weight of the penguin (grams).

$X^2 =$ the flipper length of the penguin (mm).

Label $Y \in \mathcal{Y} = \{0, 1\}$

$$Y = \begin{cases} 1 & \text{if the penguin is an Adelie} \\ 0 & \text{if the penguin is a Chinstrap.} \end{cases}$$

Example: Penguin classification

Now let's carry out a train test split.

```
num_total<-peng_total%>%nrow() # number of penguin data
num_train<-floor(num_total*0.75) # number of train examples
num_test<-num_total-num_train # number of test samples

set.seed(1) # set random seed for reproducibility
test_inds<-sample(seq(num_total),num_test) # random sample of test indices
train_inds<-setdiff(seq(num_total),test_inds) # training data indices

peng_train<-peng_total%>%filter(row_number() %in% train_inds) # train data
peng_test<-peng_total%>%filter(row_number() %in% test_inds) # test data
```

Remember to set a random seed for reproducibility.

Example: Penguin classification

Now let's carry out a train test split.

```
num_total<-peng_total%>%nrow() # number of penguin data
num_train<-floor(num_total*0.75) # number of train examples
num_test<-num_total-num_train # number of test samples

set.seed(1) # set random seed for reproducibility
test_inds<-sample(seq(num_total),num_test) # random sample of test indices
train_inds<-setdiff(seq(num_total),test_inds) # training data indices

peng_train<-peng_total%>%filter(row_number() %in% train_inds) # train data
peng_test<-peng_total%>%filter(row_number() %in% test_inds) # test data
```

We can also separate out the feature vectors and labels.

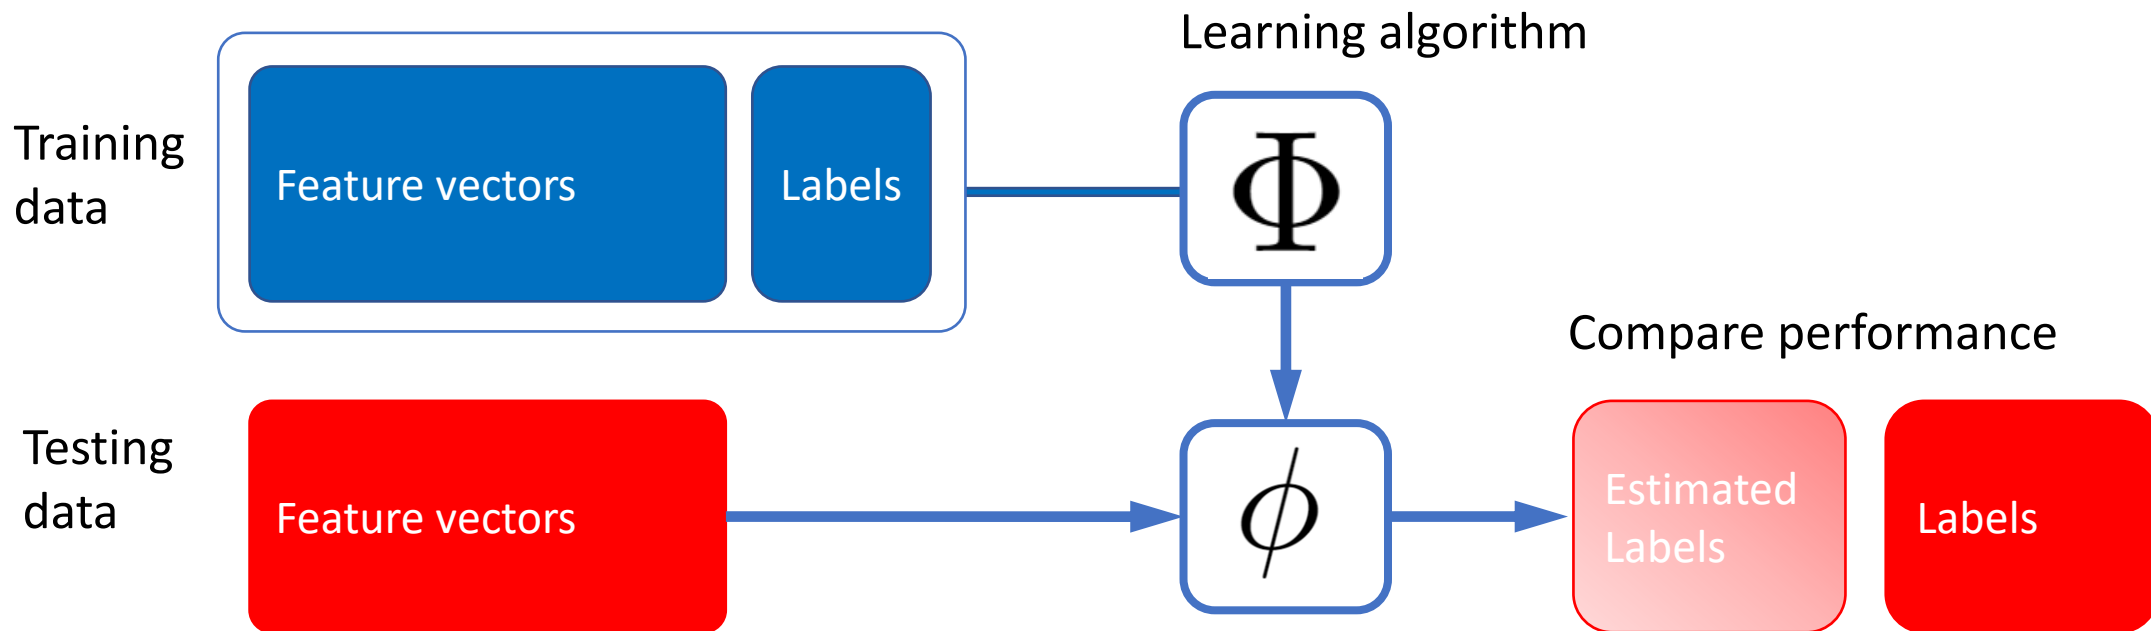
```
peng_train_x<-peng_train%>%select(-species) # train feature vectors
peng_train_y<-peng_train%>%pull(species) # train labels

peng_test_x<-peng_test%>%select(-species) # test feature vectors
peng_test_y<-peng_test%>%pull(species) # test labels
```

The train test split

Key point: Never use your test data to learn your classifier!

The goal of the test data is to see how well the classification rule does on unseen data.



Now take a break!



Statistical Computing & Empirical Methods

A probabilistic model for classification

We begin with a feature space \mathcal{X} . In the simplest case this will be $\mathcal{X} = \mathbb{R}^d$.

We have a finite set of categories \mathcal{Y} . In the simplest case this will be $\mathcal{Y} = \{0, 1\}$.

We then have random variables (X, Y)

X is a feature vector which takes values in \mathcal{X} .

Y is a label which takes values in \mathcal{Y} .

The random variables $(X, Y) \sim P$ have joint distribution P .

A probabilistic model for classification

We begin with a feature space \mathcal{X} . In the simplest case this will be $\mathcal{X} = \mathbb{R}^d$.

We have a finite set of categories \mathcal{Y} . In the simplest case this will be $\mathcal{Y} = \{0, 1\}$.

We then have random variables $(X, Y) \sim P$ with joint distribution P .

In addition we have some training data $\mathcal{D} = ((X_1, Y_1), \dots, (X_n, Y_n))$.

We assume that examples $(X_i, Y_i) \sim P$ are independent and identically distributed.

This will let us learn properties about the underlying distribution of $(X, Y) \sim P$.

Measuring performance

We have random variables $(X, Y) \sim P$ with joint distribution P .

Our goal of the learn a classification rule $\phi : \mathcal{X} \rightarrow \mathcal{Y}$, also known as classifier,

such that $\phi(X) \approx Y$ for typical $(X, Y) \sim P$.

We quantify our performance with the **expected test error**

$$\mathcal{R}(\phi) := \mathbb{P}(\phi(X) \neq Y).$$

A good classifier $\phi : \mathcal{X} \rightarrow \mathcal{Y}$ is one with a low expected test error.

Measuring performance

We quantify our performance with the **test error**

$$\mathcal{R}(\phi) := \mathbb{P}(\phi(X) \neq Y) .$$

This is the average number of mistakes on **unseen** data.

A good classifier $\phi : \mathcal{X} \rightarrow \mathcal{Y}$ is one with a low test error.

This weights all types of errors equally. Sometimes we weight different types of errors differently.

In practice we are also interested in computational issues e.g. Training/testing time and memory.

The Bayes classifier

We quantify our performance with the **test error**

$$\mathcal{R}(\phi) := \mathbb{P}(\phi(X) \neq Y).$$

The **Bayes classifier** or Bayes classification rule is the **best possible** classifier from an error perspective,

$$\mathcal{R}(\phi^*) = \min \{ \mathcal{R}(\phi) : \phi : \mathcal{X} \rightarrow \mathcal{Y} \text{ is a classifier} \}.$$

That is, $\phi^* : \mathcal{X} \rightarrow \mathcal{Y}$ is a classifier which **minimizes the test error** over all possible classifiers.

The Bayes classifier

The **Bayes classifier** $\phi^* : \mathcal{X} \rightarrow \mathcal{Y}$ minimizes the test error over all possible classifiers

$$\mathcal{R}(\phi^*) = \min \{ \mathbb{P}(\phi(X) \neq Y) : \phi : \mathcal{X} \rightarrow \mathcal{Y} \text{ is a classifier} \}.$$

Let's think about the binary case where $\mathcal{Y} = \{0, 1\}$.

We can define the Bayes classifier in terms of probability as follows,

$$\phi^*(x) := \begin{cases} 1 & \text{if } \mathbb{P}(Y = 1|X = x) \geq \mathbb{P}(Y = 0|X = x) \\ 0 & \text{if } \mathbb{P}(Y = 0|X = x) > \mathbb{P}(Y = 1|X = x). \end{cases}$$

Now take a break!



Statistical Computing & Empirical Methods

Learning from data

The **Bayes classifier** $\phi^* : \mathcal{X} \rightarrow \mathcal{Y}$ minimizes the test error over all possible classifiers

$$\mathcal{R}(\phi^*) = \min \{ \mathbb{P}(\phi(X) \neq Y) : \phi : \mathcal{X} \rightarrow \mathcal{Y} \text{ is a classifier} \}.$$

In an ideal world our computer would already know a lot about the distribution of $(X, Y) \sim P$

If this were the case, we could just use mimic the Bayes classifier $\phi^* : \mathcal{X} \rightarrow \mathcal{Y}$.

Unfortunately, the computer doesn't have prior knowledge of $(X, Y) \sim P$.

Instead we rely upon learning algorithms to learn information about the underlying distribution from the training data $\mathcal{D} = ((X_1, Y_1), \dots, (X_n, Y_n))$ i.i.d. with $(X_i, Y_i) \sim P$.

Test error vs. train error

The true goal is to find a classification rule $\phi : \mathcal{X} \rightarrow \mathcal{Y}$ with a low test error,

$$\mathcal{R}(\phi) := \mathbb{P}(\phi(X) \neq Y) \text{ with } (X, Y) \sim P.$$

Unfortunately, we can't directly observe the test error $\mathcal{R}(\phi) := \mathbb{P}(\phi(X) \neq Y)$.

However, we can use the training data $\mathcal{D} = ((X_1, Y_1), \dots, (X_n, Y_n))$

to compute the **train error**

$$\hat{\mathcal{R}}_n(\phi) := \frac{1}{n} \sum_{i=1}^n \mathbb{1} \{ \phi(X_i) \neq Y_i \}.$$

Test error vs. train error

Given a classifier $\phi : \mathcal{X} \rightarrow \mathcal{Y}$ we can't directly compute the **test error**

$\mathcal{R}(\phi) := \mathbb{P}(\phi(X) \neq Y) \approx$ Average number of mistakes on unseen data.

We can use $\mathcal{D} = ((X_1, Y_1), \dots, (X_n, Y_n))$ to compute the **train error**

$\hat{\mathcal{R}}_n(\phi) := \frac{1}{n} \sum_{i=1}^n \mathbb{1} \{\phi(X_i) \neq Y_i\} \approx$ Average number of mistakes on training data.

Key point: The train error and the test error are not the same.

However, we can use information about the training data to achieve a low-test error.

Test error vs. train error

Given $\phi : \mathcal{X} \rightarrow \mathcal{Y}$ we can't directly compute the **test error** $\mathcal{R}(\phi) := \mathbb{P}(\phi(X) \neq Y)$

We can use \mathcal{D} to compute the **train error** $\hat{\mathcal{R}}_n(\phi) := \frac{1}{n} \sum_{i=1}^n \mathbb{1} \{ \phi(X_i) \neq Y_i \}$.

Example: Image classification

Train error - average number of misclassified images in the training data.

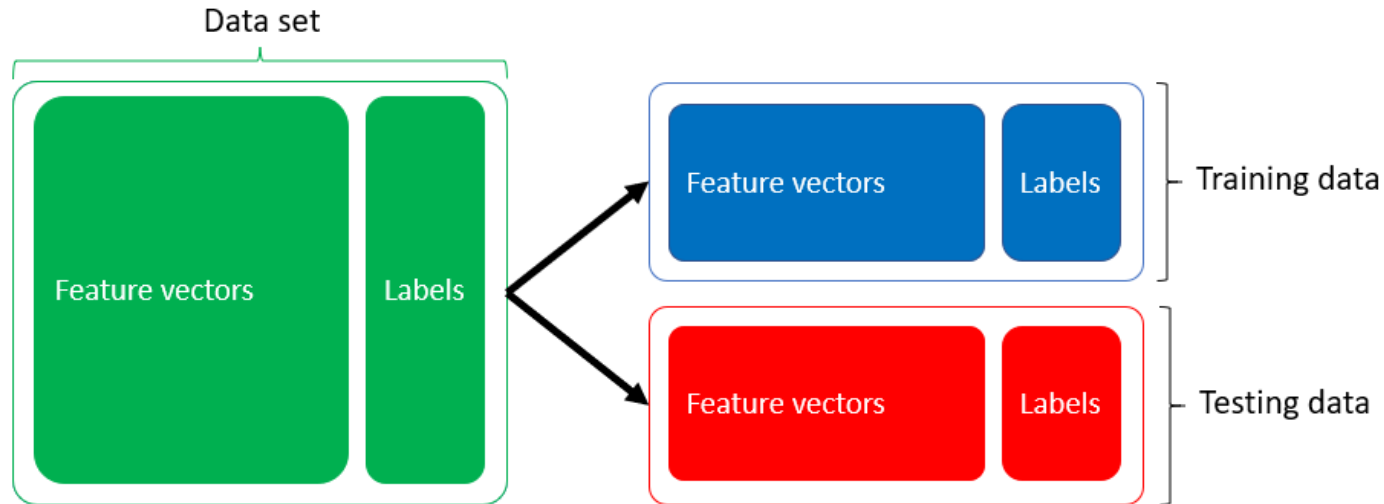
Test error – average number of misclassified future as yet unseen images.

Our real goal is to do well on the unseen data!



The train test split

For the classification rule to be successful it must perform well on **unseen** data.



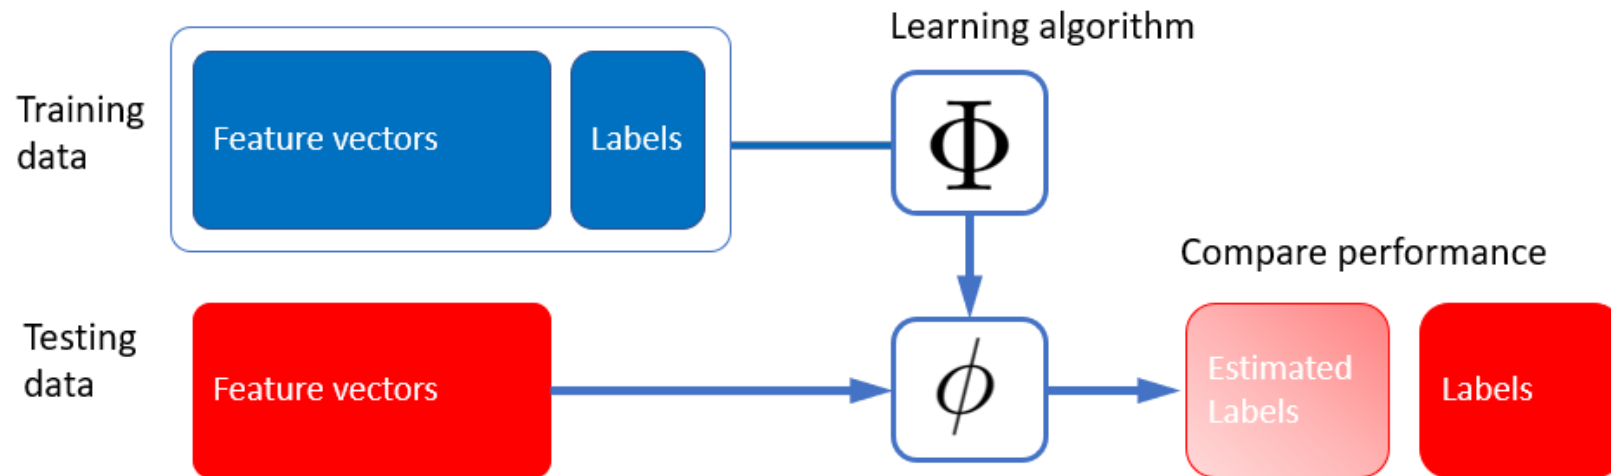
We will use the test data to estimate the test error:

$$\mathcal{R}(\phi) = \mathbb{P}(\phi(X) \neq Y) \approx \text{Average number of mistakes on unseen data.}$$

Not the same as the train error: The average number of mistakes on the training data.

The train test split

For the classification rule to be successful it must perform well on **unseen** data.



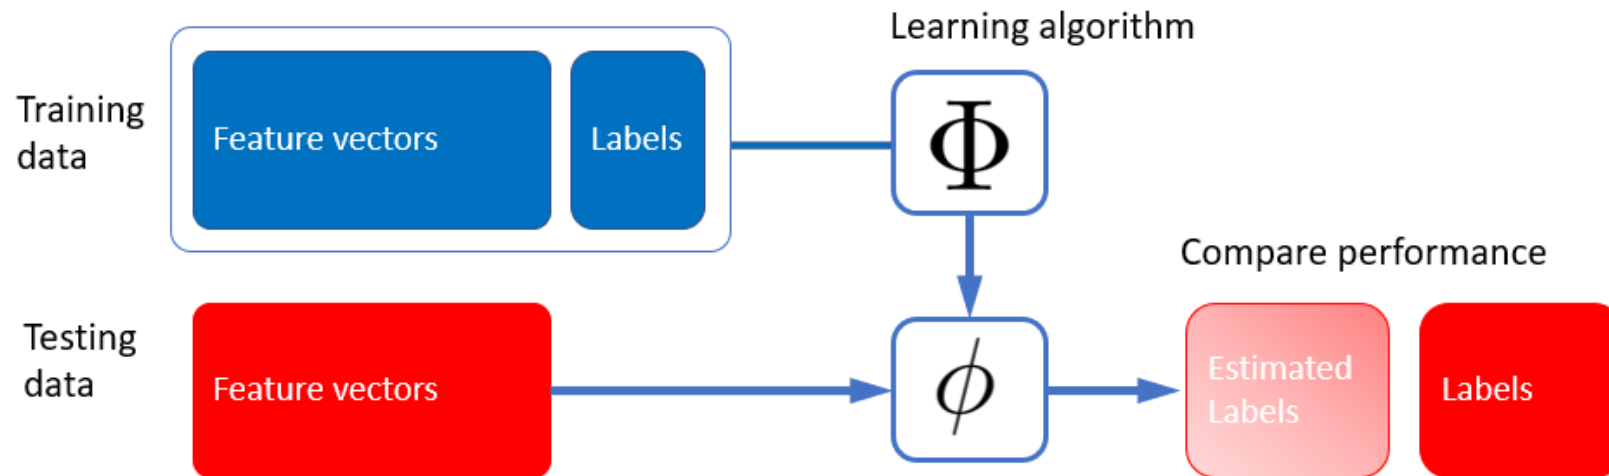
We will use the test data to estimate the test error:

$$\mathcal{R}(\phi) = \mathbb{P}(\phi(X) \neq Y) \approx \text{Average number of mistakes on unseen data.}$$

Not the same as the train error: The average number of mistakes on the training data.

Learning algorithms

For the classification rule to be successful it must perform well on **unseen** data.

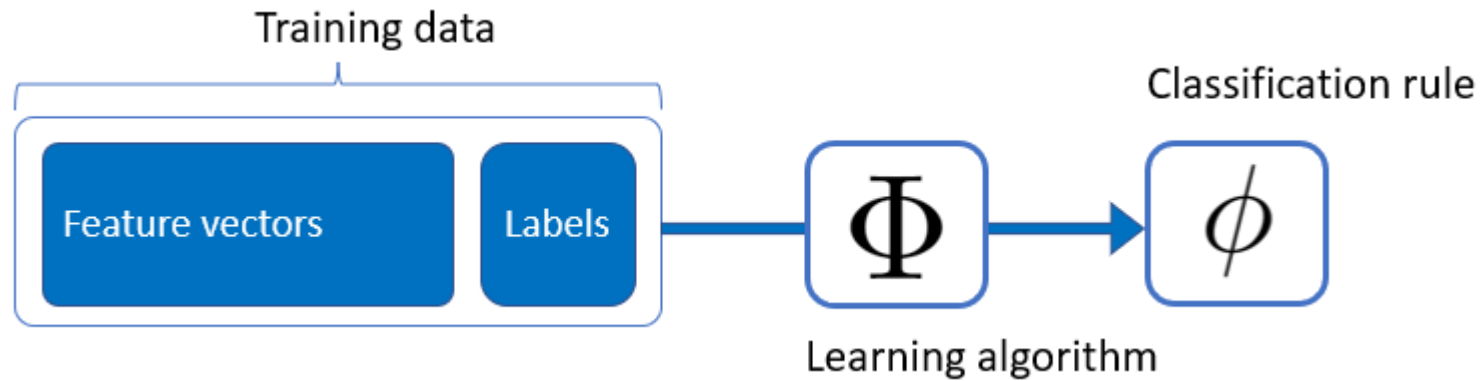


Once we have understood the classification problem.

We can start talking about how to solve the problem: Learning algorithms.

Learning algorithms

Learning algorithms are the rules for converting training data into classifiers.



Examples:

Linear Discriminant Analysis, Logistic Regression, Nearest Neighbor classifiers, Random Forests, Boosting, Neural networks, SVMs.



What have we covered today?

- We began by introducing the concept of classification with some examples.
- We emphasized the importance of predictive performance on unseen data.
- We used probabilistic ideas to understand the classification problem.
- We emphasized the difference the difference between train and test error.
- We discussed the fundamental concept of a Bayes classifier.
- We also considered the supervised learning pipeline and the role of the test-train split.



University of
BRISTOL

Thanks for listening!

henry.reeve@bristol.ac.uk

Include EMATM0061 in the subject of your email.

Statistical Computing & Empirical Methods (EMATM0061)

MSc in Data Science, Teaching block 1, 2021.

What have we covered today?

